

```
!pip install nltk
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: nltk in /usr/local/lib/python3.7/dist-
packages (3.7)
Requirement already satisfied: click in /usr/local/lib/python3.7/dist-
packages (from nltk) (7.1.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-
packages (from nltk) (4.64.1)
Requirement already satisfied: regex<=2021.8.3 in
/usr/local/lib/python3.7/dist-packages (from nltk) (2022.6.2)
Requirement already satisfied: joblib in
/usr/local/lib/python3.7/dist-packages (from nltk) (1.2.0)
```

```
df = pd.read_csv('/content/spam.csv', encoding='latin-1')
df.sample(5)
```

```
      v1      v2 Unnamed:
2  \
1667  ham  So now my dad is gonna call after he gets out ...
NaN
3556  ham          I had it already..sabarish asked me to go..
NaN
2729  spam  Urgent! Please call 09066612661 from your land...
NaN
3866  ham          Alright we're hooked up, where you guys at
NaN
2760  ham          I dont thnk its a wrong calling between us
NaN
```

```
      Unnamed: 3 Unnamed: 4
1667      NaN      NaN
3556      NaN      NaN
2729      NaN      NaN
3866      NaN      NaN
2760      NaN      NaN
```

```
df.shape
```

```
(5572, 5)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
```

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	v1	5572 non-null	object
1	v2	5572 non-null	object
2	Unnamed: 2	50 non-null	object
3	Unnamed: 3	12 non-null	object
4	Unnamed: 4	6 non-null	object

dtypes: object(5)

memory usage: 217.8+ KB

```
df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], inplace=True)
df.sample(5)
```

	v1	v2
4286	ham	I wud never mind if u dont miss me or if u don...
2103	ham	Its a site to simulate the test. It just gives...
4726	spam	Had your mobile 10 mths? Update to the latest ...
2350	ham	You will be in the place of that man
2733	ham	Do I_ noe if ben is going?

```
df.rename(columns={'v1': 'target', 'v2': 'text'}, inplace=True)
df.sample(5)
```

	target	text
5439	ham	Hey i've booked the 2 lessons on sun liao...
3303	ham	IM GONNAMISSU SO MUCH!!I WOULD SAY IL SEND U A...
1946	ham	Hey we can go jazz power yoga hip hop kb and y...
2536	ham	You do what all you like
435	ham	The message sent is askin for <#> dolla...

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
df['target'] = encoder.fit_transform(df['target'])
df.head()
```

	target	text
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
df.isnull().sum()
```

```
target    0
text      0
dtype: int64
```

```
df.duplicated().sum()
```

403

```

df = df.drop_duplicates(keep='first')
df.duplicated().sum()

0

df.shape

(5169, 2)


df.head()
   target      text
0       0  Go until jurong point, crazy.. Available only ...
1       0           Ok lar... Joking wif u oni...
2       1  Free entry in 2 a wkly comp to win FA Cup fina...
3       0  U dun say so early hor... U c already then say...
4       0  Nah I don't think he goes to usf, he lives aro...

df['target'].value_counts()

0    4516
1     653
Name: target, dtype: int64

plt.pie(df['target'].value_counts(),
labels=['ham', 'spam'],autopct="%0.2f")
plt.show()

```



```

import nltk
nltk.download('punkt')

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.

```

True

```
df['num_characters'] = df['text'].apply(len)
df.head()
```

	target	text
num_characters		
0	0	Go until jurong point, crazy.. Available only ...
111		
1	0	Ok lar... Joking wif u oni...
29		
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
155		
3	0	U dun say so early hor... U c already then say...
49		
4	0	Nah I don't think he goes to usf, he lives aro...
61		

```
df['num_words'] = df['text'].apply(lambda
x:len(nltk.word_tokenize(x)))
df.head()
```

	target	text
num_characters \		
0	0	Go until jurong point, crazy.. Available only ...
111		
1	0	Ok lar... Joking wif u oni...
29		
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
155		
3	0	U dun say so early hor... U c already then say...
49		
4	0	Nah I don't think he goes to usf, he lives aro...
61		

	num_words
0	24
1	8
2	37
3	13
4	15

```
df['num_sentences'] = df['text'].apply(lambda
x:len(nltk.sent_tokenize(x)))
df.head()
```

	target	text
num_characters \		
0	0	Go until jurong point, crazy.. Available only ...
111		
1	0	Ok lar... Joking wif u oni...

```

29
2      1  Free entry in 2 a wkly comp to win FA Cup fina...
155
3      0  U dun say so early hor... U c already then say...
49
4      0  Nah I don't think he goes to usf, he lives aro...
61

```

	num_words	num_sentences
0	24	2
1	8	2
2	37	2
3	13	1
4	15	1

```
df[['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences
count	5169.000000	5169.000000	5169.000000
mean	78.977945	18.453279	1.947185
std	58.236293	13.324793	1.362406
min	2.000000	1.000000	1.000000
25%	36.000000	9.000000	1.000000
50%	60.000000	15.000000	1.000000
75%	117.000000	26.000000	2.000000
max	910.000000	220.000000	28.000000

```
df[df['target'] == 0]
```

```
[['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences
count	4516.000000	4516.000000	4516.000000
mean	70.459256	17.120903	1.799601
std	56.358207	13.493725	1.278465
min	2.000000	1.000000	1.000000
25%	34.000000	8.000000	1.000000
50%	52.000000	13.000000	1.000000
75%	90.000000	22.000000	2.000000
max	910.000000	220.000000	28.000000

```
df[df['target'] == 1]
```


```
[['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences
count	653.000000	653.000000	653.000000
mean	137.891271	27.667688	2.967841
std	30.137753	7.008418	1.483201
min	13.000000	2.000000	1.000000
25%	132.000000	25.000000	2.000000
50%	149.000000	29.000000	3.000000

```
75%          157.000000    32.000000         4.000000
max          224.000000    46.000000         8.000000
```


```
plt.figure(figsize=(12,6))
sns.histplot(df[df['target'] == 0]['num_characters'], color='green')
sns.histplot(df[df['target'] == 1]['num_characters'], color='red')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fe65436d790>



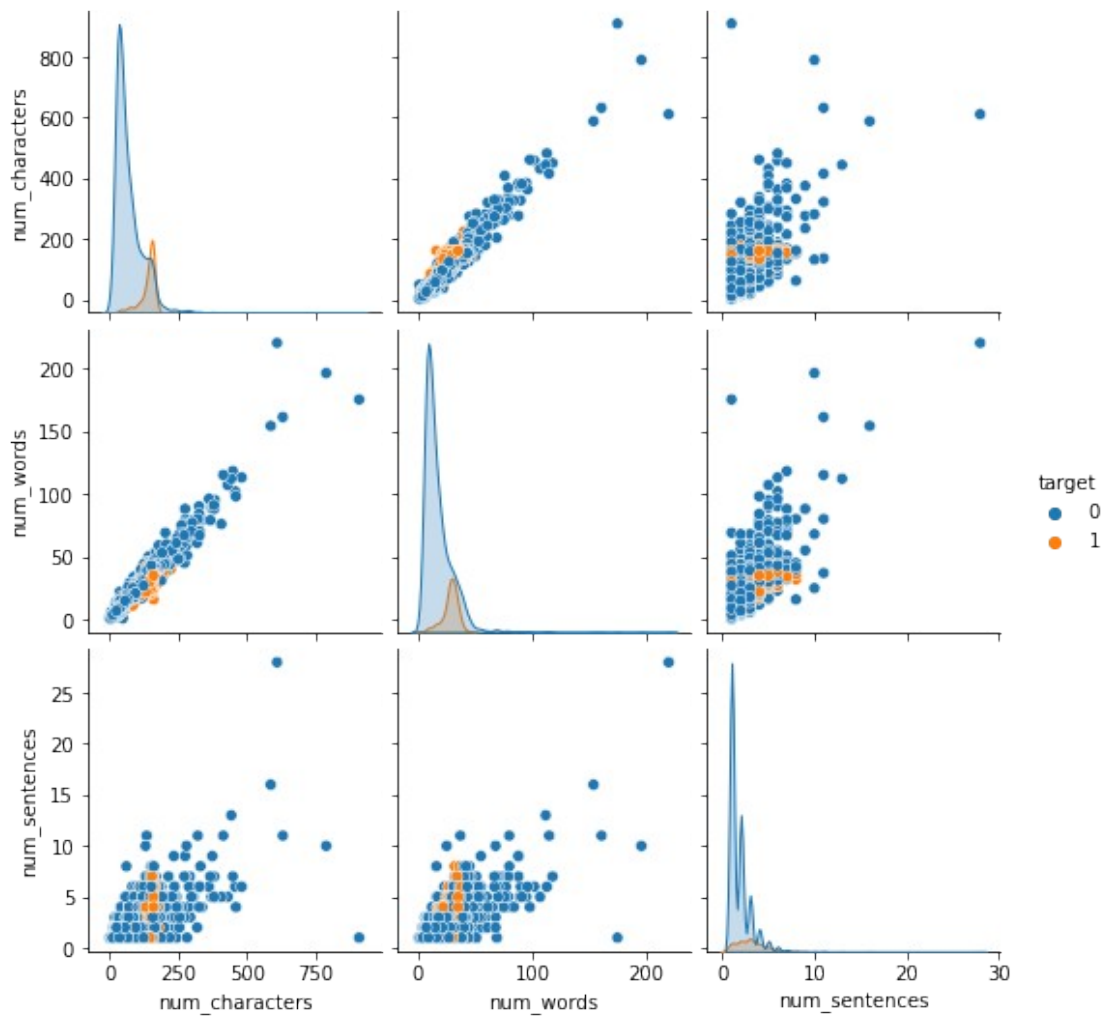
```
plt.figure(figsize=(12,6))
sns.histplot(df[df['target'] == 0]['num_words'], color='green')
sns.histplot(df[df['target'] == 1]['num_words'], color='red')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fe6540c56d0>



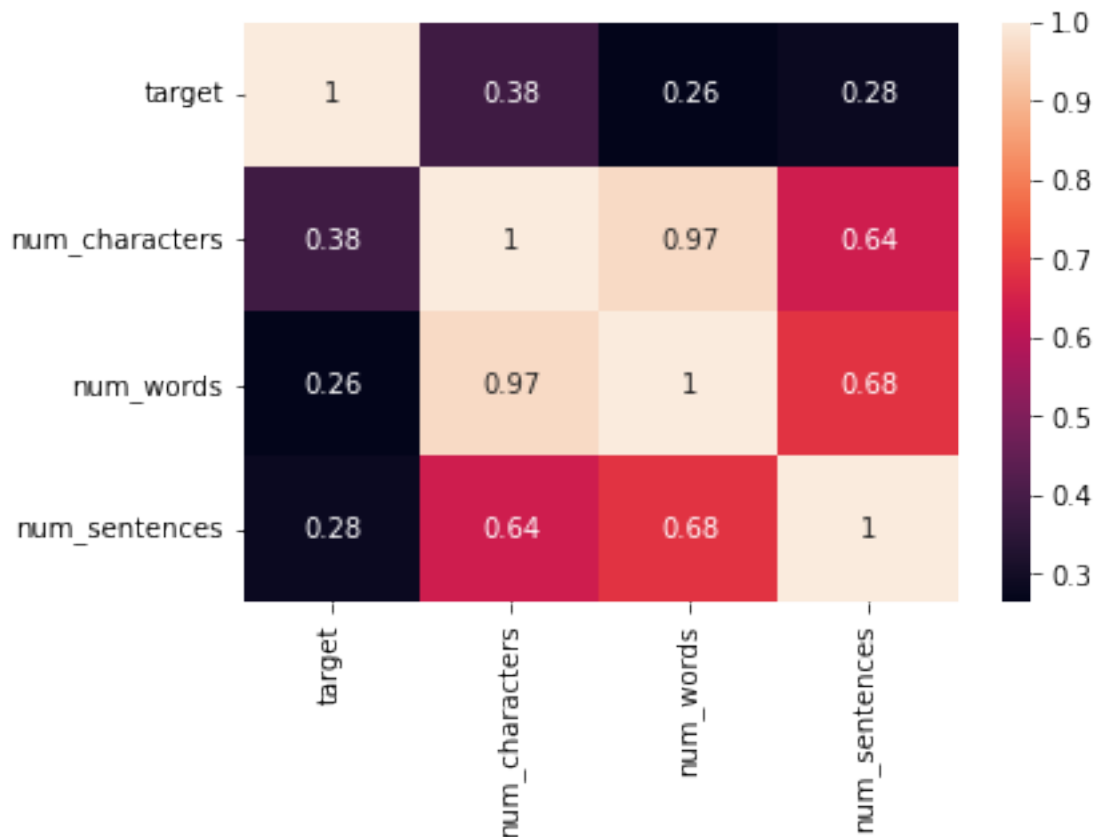
```
sns.pairplot(df, hue='target')
```

<seaborn.axisgrid.PairGrid at 0x7fe653e9a250>



```
sns.heatmap(df.corr(), annot=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fe653fd6410>



```
from nltk.stem.porter import PorterStemmer
ps = PorterStemmer()
print(ps.stem('played'))
print(ps.stem('playing'))
```

```
play
play
```

```
import string
string.punctuation
```

```
{"type": "string"}
```

```
nltk.download('stopwords')
from nltk.corpus import stopwords
stopwords.words('english')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```

```
['i',
 'me',
 'my',
 'myself',
 'we',
 'our',
```


'ours',
'ourselves',
'you',
"you're",
"you've",
"you'll",
"you'd",
'your',
'yours',
'yourself',
'yourselves',
'he',
'him',
'his',
'himself',
'she',
"she's",
'her',
'hers',
'herself',
'it',
"it's",
'its',
'itself',
'they',
'them',
'their',
'theirs',
'themselves',
'what',
'which',
'who',
'whom',
'this',
'that',
"that'll",
'these',
'those',
'am',
'is',
'are',
'was',
'were',
'be',
'been',
'being',
'have',
'has',
'had',
'having',

'do',
'does',
'did',
'doing',
'a',
'an',
'the',
'and',
'but',
'if',
'or',
'because',
'as',
'until',
'while',
'of',
'at',
'by',
'for',
'with',
'about',
'against',
'between',
'into',
'through',
'during',
'before',
'after',
'above',
'below',
'to',
'from',
'up',
'down',
'in',
'out',
'on',
'off',
'over',
'under',
'again',
'further',
'then',
'once',
'here',
'there',
'when',
'where',
'why',
'how',

'all',
'any',
'both',
'each',
'few',
'more',
'most',
'other',
'some',
'such',
'no',
'nor',
'not',
'only',
'own',
'same',
'so',
'than',
'too',
'very',
's',
't',
'can',
'will',
'just',
'don',
"don't",
'should',
"should've",
'now',
'd',
'll',
'm',
'o',
're',
've',
'y',
'ain',
'aren',
"aren't",
'couldn',
"couldn't",
'didn',
"didn't",
'doesn',
"doesn't",
'hadn',
"hadn't",
'hasn',
"hasn't",

```
'haven',  
"haven't",  
'isn',  
"isn't",  
'ma',  
'mightn',  
"mightn't",  
'mustn',  
"mustn't",  
'needn',  
"needn't",  
'shan',  
"shan't",  
'shouldn',  
"shouldn't",  
'wasn',  
"wasn't",  
'weren',  
"weren't",  
'won',  
"won't",  
'wouldn',  
"wouldn't"]
```

```
def transform_text(text):  
    text = text.lower()  
    text = nltk.word_tokenize(text)  
  
    y = []  
    for i in text:  
        if i.isalnum():  
            y.append(i)  
  
    text = y[:]  
    y.clear()  
  
    for i in text:  
        if i not in stopwords.words('english') and i not in  
string.punctuation:  
            y.append(i)  
  
    text = y[:]  
    y.clear()  
  
    for i in text:  
        y.append(ps.stem(i))  
  
    return " ".join(y)  
  
df['text'][23]
```

```

{"type": "string"}

transform_text(df['text'][23])

{"type": "string"}

df['transformed_text'] = df['text'].apply(transform_text)
df.head()

    target                                text
num_characters \
0          0  Go until jurong point, crazy.. Available only ...
111
1          0                                Ok lar... Joking wif u oni...
29
2          1  Free entry in 2 a wkly comp to win FA Cup fina...
155
3          0  U dun say so early hor... U c already then say...
49
4          0  Nah I don't think he goes to usf, he lives aro...
61

    num_words  num_sentences
transformed_text
0          24              2  go jurong point crazi avail bugi n great
world...
1           8              2                                ok lar joke
wif u oni
2          37              2  free entri 2 wkli comp win fa cup final
tkit 21...
3          13              1                                u dun say earli hor u c
alreadi say
4          15              1                                nah think goe usf live
around though

from wordcloud import WordCloud
wc =
WordCloud(width=500,height=500,min_font_size=10,background_color='white')

spam_corpus = []
for msg in df[df['target'] == 1]['transformed_text'].tolist():
    for word in msg.split():
        spam_corpus.append(word)
len(spam_corpus)


9939

from collections import Counter
from collections import Counter
sns.barplot(pd.DataFrame(Counter(spam_corpus).most_common(30))
[0],pd.DataFrame(Counter(spam_corpus).most_common(30))[1])

```

```
plt.xticks(rotation='vertical')
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.
FutureWarning




```
ham_corpus = []
for msg in df[df['target'] == 0]['transformed_text'].tolist():
    for word in msg.split():
        ham_corpus.append(word)
len(ham_corpus)
```

35394

```
from collections import Counter
sns.barplot(pd.DataFrame(Counter(ham_corpus).most_common(30))
[0],pd.DataFrame(Counter(ham_corpus).most_common(30))[1])
plt.xticks(rotation='vertical')
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in

an error or misinterpretation.
FutureWarning



```
from sklearn.feature_extraction.text import TfidfVectorizer

tfidf = TfidfVectorizer(max_features=3000) # 3000 is giving good
accuracy and precision

X = tfidf.fit_transform(df['transformed_text']).toarray()
X.shape

(5169, 3000)

y = df['target'].values

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.2,random_state=2)

from sklearn.naive_bayes import GaussianNB,MultinomialNB,BernoulliNB
from sklearn.metrics import
accuracy_score,confusion_matrix,precision_score

gnb = GaussianNB()
mnb = MultinomialNB()
bnb = BernoulliNB()
gnb.fit(X_train,y_train)
```

```

y_pred1 = gnb.predict(X_test)

print(accuracy_score(y_test,y_pred1))
print(confusion_matrix(y_test,y_pred1))
print(precision_score(y_test,y_pred1))

0.8694390715667312
[[788 108]
 [ 27 111]]
0.5068493150684932

mnb.fit(X_train,y_train)
y_pred2 = mnb.predict(X_test)
print(accuracy_score(y_test,y_pred2))
print(confusion_matrix(y_test,y_pred2))
print(precision_score(y_test,y_pred2))

0.9709864603481625
[[896   0]
 [ 30 108]]
1.0

bnb.fit(X_train,y_train)
y_pred3 = bnb.predict(X_test)
print(accuracy_score(y_test,y_pred3))
print(confusion_matrix(y_test,y_pred3))
print(precision_score(y_test,y_pred3))

0.9835589941972921
[[895   1]
 [ 16 122]]
0.991869918699187

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier

svc = SVC(kernel='sigmoid', gamma=1.0)
knc = KNeighborsClassifier()
mnb = MultinomialNB()
dtc = DecisionTreeClassifier(max_depth=5)

```



```

lrc = LogisticRegression(solver='liblinear', penalty='l1')
rfc = RandomForestClassifier(n_estimators=50, random_state=2)
abc = AdaBoostClassifier(n_estimators=50, random_state=2)
bc = BaggingClassifier(n_estimators=50, random_state=2)
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
gbdt = GradientBoostingClassifier(n_estimators=50, random_state=2)
xgb = XGBClassifier(n_estimators=50, random_state=2)

```

```

clfs = {
    'SVC' : svc,
    'KN' : knn,
    'NB' : nb,
    'DT' : dtc,
    'LR' : lrc,
    'RF' : rfc,
    'AdaBoost' : abc,
    'BgC' : bc,
    'ETC' : etc,
    'GBDT' : gbdt,
    'xgb' : xgb
}

```

```

def train_classifier(clf, X_train, y_train, X_test, y_test):
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)

    return accuracy, precision

```

```

accuracy_scores = []
precision_scores = []

```

```

for name, clf in clfs.items():

```

```

    current_accuracy, current_precision = train_classifier(clf,
X_train, y_train, X_test, y_test)

```

```

    print("For ", name)
    print("Accuracy - ", current_accuracy)
    print("Precision - ", current_precision)

```

```

    accuracy_scores.append(current_accuracy)
    precision_scores.append(current_precision)

```

```

For SVC
Accuracy - 0.9758220502901354
Precision - 0.9747899159663865
For KN

```

```

Accuracy - 0.9052224371373307
Precision - 1.0
For NB
Accuracy - 0.9709864603481625
Precision - 1.0
For DT
Accuracy - 0.9284332688588007
Precision - 0.82
For LR
Accuracy - 0.9584139264990329
Precision - 0.9702970297029703
For RF
Accuracy - 0.9748549323017408
Precision - 0.9827586206896551
For AdaBoost
Accuracy - 0.960348162475822
Precision - 0.9292035398230089
For BgC
Accuracy - 0.9574468085106383
Precision - 0.8671875
For ETC
Accuracy - 0.9748549323017408
Precision - 0.9745762711864406
For GBDT
Accuracy - 0.9477756286266924
Precision - 0.92
For xgb
Accuracy - 0.9439071566731141
Precision - 0.9347826086956522

```

```

df =
pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy':accuracy_scores,'Precision':precision_scores}).sort_values('Precision',ascending=False)
df

```

	Algorithm	Accuracy	Precision
1	KN	0.905222	1.000000
2	NB	0.970986	1.000000
5	RF	0.974855	0.982759
0	SVC	0.975822	0.974790
8	ETC	0.974855	0.974576
4	LR	0.958414	0.970297
10	xgb	0.943907	0.934783
6	AdaBoost	0.960348	0.929204
9	GBDT	0.947776	0.920000
7	BgC	0.957447	0.867188
3	DT	0.928433	0.820000